

Venkata Shiva Krishna Reddy Avula

Email- [vavula1@uncc.edu](mailto:vavula1@uncc.edu)

Id- 800963515

1. A. Please find the code heat.cpp file in the rar.

The partitioning was horizontal partitioning among nodes.

B.and C.

Strong Scaling Experiment      Procs      Time in Seconds

1GB	1	17.86
	2	10.92
	4	5.5
	8	2.76
	16	1.69
	24	1.10
4Gb	1	65
	2	41.28
	4	20.73
	8	10.81
	16	6.24
	24	6.48
10Gb	1	-
	2	107.35
	4	53.9
	8	27.7
	16	16.6
	24	10.89
80	1	-
	2	-
	4	-
	8	137.99
	16	82.5
	24	54

Weak Scaling Experiment      Procs      Time in seconds

500	1	8.17
	2	10.58
	4	10.65
	8	10.85
	16	12.7
	24	12.93
1gb	1	16.9
	2	21.55
	4	23.99
	8	22.33
	16	26.44
	24	26.65
2gb	1	33.04
	2	42.07
	4	41.91
	8	43.07


Venkata Shiva Krishna Reddy Avula

Email- [vavula1@uncc.edu](mailto:vavula1@uncc.edu)


Id- 800963515

	16	51
	24	51.86

Graphs for weak Scaling




Graphs for Strong Scaling



D.

2. A. Please find the matmul.cpp code file

D.



D. The computation and communication can be overlapped by first calculating the elements of row succeeding the 1<sup>st</sup> data row in each node till the second last row. In the meanwhile, the nodes can send and receive data between the adjacent nodes through MPI\_ISEND and MPI\_IRECV. Before calculating the first and last data row each node can wait to check whether all the communication has happened and if it finishes they can calculate the first and last rows in each node in each iteration.

2. A. Please consider the mat\_mul.cpp code file.

Initializing the X array on top row nodes>>>

Initializing the A array in each node.>>

Iterations =0

While(iterations <20){

Sending the X array to the nodes in same column>>

Nodes that are not in the first receive the X array>>

Calculating the partial sum in the X array>>

Row wise reduction of elements on the nodes in first column>>

Sending the reduced sum X array to the nodes in top row.>>

The top row nodes except node0 receive the X array

Iterations++}

D. The communication and computation overlap can be done for the large matrices on each node by further dividing the Xarray. So that when a part of Xarray is received each node start calculation for the corresponding indices and store at the first index. Meanwhile, second part of X array can be received and so on. And for the reduction we could first reduce on each node and the reduction can be done on all nodes before sending it to the top row nodes for each iteration.